

Yijiang (River) Dong

Phone: (+44) 7548809109 | Email: yd358@cam.ac.uk | [Google Scholar](#)

Education

University of Cambridge; Cambridge, UK

January 2024—Present

- PhD in Natural Language Processing
- Advisors: Nigel Collier, Ivan Vulic
- Research Interests: Large Language Models, Personalization, Alignment, RLHF

University of Pennsylvania; Philadelphia, US

September 2019—September 2023

- BA in Mathematics; MSE in Data Science; GPA: 3.9/4.0
- Advisors: Chris Callison-Burch, Lyle Ungar

Publications

Can LLM be a Personalized Judge? [\[paper\]](#)

[Yijiang River Dong*](#), [Tiancheng Hu*](#), Nigel Collier

In Findings of Empirical Methods in Natural Language Processing (EMNLP 2024)

Unmemorization in Large Language Models via Self-Distillation and Deliberate Imagination [\[paper\]](#)

[Yijiang River Dong](#), Hongzhou Lin, Mikhail Belkin, Ramon Huerta, Ivan Vulic

In Submission for NAACL 2025

CORRPUS: Codex-Leveraged Structured Representations for Neuro-symbolic Story Understanding [\[paper\]](#)

[Yijiang River Dong](#), Lara J Martin, Chris Callison-Burch

In Findings of the Annual Meeting of the Association for Computational Linguistics (ACL 2023)

COTTAGE: Coherent Text Adventure Games Generation [\[paper\]](#)

[Yijiang River Dong](#)

Master's Thesis

Industry Experience

Machine Learning Engineer Intern, Roblox

May 2022—August 2022

- Worked in the Bots and Spam team to design text mining algorithms for bot detection
- Purposed and implemented recursive text mining algorithm and statistical test methods to detect abnormal username patterns using Spark, Hive, SQL, Hadoop, Airflow
- Improved the detection coverage by 600% while maintaining F1 score of 90%
- Implemented a Markov chain-based gibberish detector

Applied Scientist Intern, Amazon

Jan 2022—May 2022

- Worked on Amazon Help Search to expand English information retrieval (IR) model for multi-lingual IR
- Built large scale data pipeline including scraping, annotation collection, data cleaning, data augmentation
- Developed Multilingual BM25, BERT-based bi-encoder and cross-encoder model
- Improved MRR from 22% to 45% in the same runtime by using knowledge distillation on mBERT model

Investment Intern, Sequoia Capital China

Sep 2023—Dec 2023

- Worked in Seed-stage investment team to evaluate the potential of emerging AI startups, focusing on sectors such as video generation (e.g., Pika), image editing (e.g., FotorAI), and AI security (e.g., RealAI, Nexusflow).
- Investigated 20+ startups through engaging with beta products, analyzing Key Performance Indicators for startup Growth, reading cutting-edge research papers to understand technical moat, and interviewing startup founders and industry experts